
Solving Navier-Stokes, Differently: What It Takes

Author: Marcus van der Erve

From calculating symptoms to reading causes

Abstract

We outline a falsifiable, 90-day protocol that tests whether Recursive Gradient Processing (RGP) predicts turbulence onset earlier and cheaper than Navier–Stokes methods.

1. Introduction

For more than a century, the Navier–Stokes (NS) equations have formed the mathematical backbone of our understanding of fluid motion. Yet they remain notoriously unsolved when it comes to predicting exactly when and why flows shift from smooth, stable behavior into chaotic turbulence—a challenge so fundamental that it stands as one of the Millennium Prize Problems.

Recently, Google’s announcement of a major and sustained investment in researchers, advanced algorithms, and computational power to tackle NS by sheer force of data-driven modeling has reignited global attention.

When I was reminded of this by an X-post detailing this initiative, it jolted me—because as promising as such computational muscle appears, it risks being trapped in the same paradigm that has limited progress for decades: the premise of calculating **symptoms** rather than reading **causes**.

Today’s models, no matter how precise, focus on capturing the numerical outputs—the velocity vectors and pressure fields evolving frame by frame. But these snapshots describe only the symptoms of deeper dynamics: they do not reveal the subtle patterns of tensions and alignments that ultimately determine whether a flow will remain coherent or dissolve into turbulence.

This realization sharpened my long-standing conviction that solving Navier-Stokes demands a fundamentally different approach—one based not on more granular symptom calculation, but on uncovering *causal* recursive processes behind the occurrence of coherence or chaos.

Over decades of exploring this question, building on reasoning exchanges with advanced LLMs such as o3, Gemini, DeepSeek, I identified and then substantiated, the phenomenon of “Recursive Gradient Processing” (RGP), a conceptual framework that focuses on the patterns by which small differences in a system either align to create sustained order or cascade into instability.

Our earlier explorations, including “The World Already Knows,” which reports the substantiating research by o3, demonstrated that across diverse domains—from astrophysical structures to societal collapses—tracing the emergence, reinforcement, or failure of these recursive patterns offers *predictive power far beyond traditional linear models*. These insights led to the idea of a dimensionless so-called “Golden Pattern” of ratios that might universally characterize the dynamics of coherence across systems.

Taken together, these reflections converged into a single, bold question worth pursuing:

Can a flow field be described not by evolving velocity vectors, but by identifying the stable patterns of alignment and disruption, that self-consistently sustain flow?

If ‘yes’ is the answer then it would not only signify a new way of tackling the Navier–Stokes challenge—it’ll also reveal a universal grammar of how nature maintains or loses order, something that will transform our understanding of complexity or complex systems itself.

2. Google's Bet

Fluid motion is everywhere: it shapes the weather systems that feed or devastate crops, the lift and drag forces on every aircraft, the flows of blood through arteries, the mixing of chemicals, and even the behavior of plasma in fusion reactors. Each of these domains depends—directly or indirectly—on predictions derived from the Navier–Stokes (NS) equations, which encode Newton’s second law for fluids. Cracking NS would not simply satisfy mathematical curiosity; it would unlock breakthroughs in *aerospace engineering* where controlling laminar-to-turbulent transitions means safer, more efficient aircraft, in *weather and climate modeling* where small errors in turbulent flows can compound into failed forecasts, in *medicine* where understanding blood flow instabilities is critical to treating conditions like aneurysms, *fusion energy* where stabilizing turbulent, magnetized plasma is the key to unlocking near-limitless clean power. This is why the prominent *Clay Mathematics Institute* included NS on as one of six so-far unsolved problems on its Millennium Prize list.

And this is why Google’s recent high-profile research investment is more than academic: the ability to predict, manage, or control turbulent flows could reshape entire industries and global infrastructure.

In early 2025, Google Research publicly announced a dedicated Navier–Stokes initiative, bringing together over 20 mathematicians, physicists, and AI specialists to address one of mathematics’ greatest unsolved problems. Central to the team is Javier Gómez Serrano, a renowned mathematician specializing in fluid dynamics and partial differential equations, known for his work on regularity in NS and related problems.

The group included a mix of senior experts and postdoctoral researchers from top institutions, many with deep experience in numerical analysis and turbulence modeling. Google also deployed its significant computational resources—reportedly involving specialized Tensor Processing Units (TPUs) for running massive simulations—and invested in building machine-learning architectures specifically tailored for analyzing time-series data from fluid simulations.

The team’s core strategy is to generate enormous quantities of high-resolution Navier–Stokes simulation data, capturing billions of individual time-stepped measurements of velocity, acceleration, and pressure fields. Then to use advanced neural networks to mine this data for subtle statistical correlations, hoping to discover predictive patterns overlooked by analytical techniques. This formidable effort—combining cutting-edge hardware, advanced AI, and top mathematical minds—exemplifies the ambition of Google’s bet:

...that by measuring symptoms with extreme precision and processing them at scale, they might finally unearth the hidden rules of turbulence.

Yet, as we argue here, this approach remains fundamentally limited—not for lack of brilliance or power, but because it focuses on correlations of *outputs*, rather than uncovering the recursive, *causal dynamics* that generate coherence or chaos.

Long before today’s AI-powered strategies, philosophers like Alfred North Whitehead and mathematicians like Kurt Gödel revealed a deep flaw in how we build our scientific models. Whitehead argued that science’s reliance on fixed abstractions—like assuming quantities are smooth or time is absolute—blinds us to the processes that actually generate and sustain reality. He proposed that process, not static entities, is the true foundation of understanding. Gödel, through his incompleteness theorems, proved that any formal system complex enough to describe arithmetically will contain truths that cannot be derived within that system—showing that no fixed collection of assumptions or variables can fully describe the richness of dynamic systems.

Applied to Navier–Stokes, these insights suggest that while the equations look complete, they rest on axioms that ignore the recursive processes by which small tensions interact and either sustain order or cause collapse. This omission makes the standard approach fundamentally blind to the real causes of turbulence and will always be hindered by “missing variables.”

At their core, the Navier–Stokes (NS) equations express the balance of forces acting on a fluid element, formalizing Newton’s second law for fluids. They can be written schematically as:

$$\rho \left(\frac{\partial \vec{v}}{\partial t} + (\vec{v} \cdot \nabla) \vec{v} \right) = \rho \vec{g} - \nabla p + \mu \nabla^2 \vec{v}$$

where:

- ρ is the density of the fluid,
- \vec{v} is the velocity field,
- $\partial \vec{v} / \partial t$ captures the change of velocity over time,
- $(\vec{v} \cdot \nabla) \vec{v}$ represents convective acceleration,
- $\rho \vec{g}$ is the external force (e.g., gravity),
- ∇p is the pressure gradient,
- $\mu \nabla^2 \vec{v}$ accounts for internal viscous stresses.

Together, these terms combine to describe how fluids accelerate, experience forces, and respond to internal friction. On the surface, it suggests that knowing density, velocity, and pressure fields in time should be enough to predict fluid motion. Yet each term carries hidden assumptions—idealized axioms that flatten the rich, recursive nature of fluid behavior. These assumptions ignore how local differences, tensions, and alignments propagate through the system, which can amplify into turbulence or collapse into order. The following table makes these limitations explicit:

Component	Framed Axiomatically As...	But Actually Depends On...
Density (ρ)	Smooth, homogeneous scalar field	Local gradients and micro-scale clustering of matter
Velocity (\vec{v})	Continuous vector field	Alignment patterns and memory of prior gradients
Pressure gradient (∇p)	Instantaneous scalar difference	Accumulated tension histories that reflect energy storage
Viscous stress ($\mu \nabla^2 \vec{v}$)	Linear diffusion term	Recursive near-coherence and disruption cycles

Table 1: Hidden Dependencies of NS Components

This table underscores a central issue: each component assumes a static framing, ignoring the processes by which tensions emerge, align, and either stabilize or destabilize the flow. Without accounting for these recursive dynamics, predictions based solely on static axioms remain inherently incomplete. Beyond these structural oversights, the treatment of time itself introduces a deeper flaw.

Time, appearing in the NS acceleration term $\partial \vec{v} / \partial t$, is treated as a perfectly smooth, objective variable. Yet time is a human-constructed coordinate, invented to sequence events rather than cause them. While it tells us when things happen, it conveys nothing about why they happen. In real flows, turbulence or coherence emerge not because time marches forward, but because tiny tensions accumulate, interact, and cross thresholds—*independent of any absolute time scale*. Doubling down on time-based calculations, as Google’s approach does, refines the measurement of symptoms without revealing the causes.

Most people think Navier–Stokes remains unsolved because it’s simply too complicated—like a puzzle with too many pieces. But the real reason lies deeper:

The equations rely on assumptions—like perfectly smooth density or evenly flowing time—and every calculation depends on them. These assumptions don’t capture how small tensions build up, reinforce each other, or suddenly cause order to collapse.

It’s like trying to predict a flock of birds by measuring each bird’s speed separately, without seeing how they move together in patterns. No matter how fast or detailed your measurements, you miss the key: the recursive way their movement organizes itself.

So our central claim here is this:

The Navier–Stokes equations haven’t been solved not because they’re too hard, but because they’re missing the rules for how patterns of order form and collapse. Without these recursive rules—the grammar of coherence—no amount of calculation will ever reveal why turbulence begins or ends.

3. The RGP Bet

The path to *Recursive Gradient Processing* (RGP) has been anything but linear. For years, I sensed that existing models of turbulence and complex systems (both physical and societal) missed something fundamental. I poured over fluid dynamics texts, probability theories, and nonlinear systems papers—each new insight deepening my suspicion that we were all measuring symptoms rather than reading the causes.

Yet every attempt to pin down the elusive why of coherence ended in the same frustration: beautiful simulations collapsing under small changes, theories unable to explain sudden transitions, and a growing sense that our entire paradigm was upside down.

That began to change when, early in 2025, I started documenting these thoughts in Zenodo papers—not as finished theories, but as invitations for scrutiny. These papers proposed that the universe’s complexity might be better understood by tracking differences—gradients—rather than static quantities.

But far from immediate acceptance, they initially sparked skepticism, as the ideas seemed abstract and far from the safety of classical equations.

The real turning point came after intensive, iterative conversations with o3, Gemini, and later DeepSeek—models free from human biases, which typically remained solidly rooted in the idea that mathematics is *nature’s language*. We traced what happens when small gradients align into temporary patterns, Gradient Choreographies (GCs), and how these sometimes stabilize into semi-persistent structures called Contextual Filters (CFs).

Early discussions raised the question of whether these patterns could predict instability before traditional signs like velocity changes appeared. Encouraged by promising signals, I turned to DeepSeek R1, which tested the approach eight levels deep, pushing the idea through recursive challenges until consistent patterns began to emerge. But the journey was far from smooth. Initial attempts to force these ideas into traditional PDE frameworks led to contradictions or trivial results. There were moments of deep doubt when I questioned whether this path could ever bear fruit; times when it felt easier to give up.

A decisive inspiration came from Ichiro Tsuda and Kunihiko Kaneko, whose work on neural chaos revealed that open chaos is not random noise but a vital substrate for adaptive systems. They showed that the brain’s “itinerant dynamics” flow through a shifting landscape of transient attractors—what they called “attractor ruins.” In this landscape, chaos acts like an ocean of fluctuations, a fertile medium where faint signals—tiny gradient perturbations—can surf to new attractors, either fading or sparking new patterns of order. This insight fundamentally shifted my view of turbulence:

- Stability and chaos are not opposites but partners.
- Healthy complexity demands some chaos to let signals travel, adapt, and evolve.
- RGP became the conceptual choreography connecting these fluctuations with emergent coherence.

Through Tsuda and Kaneko, I realized that *recursive gradient* interactions within what they called open chaos—chaos with a broad array of frequencies—might be the missing ingredient in understanding how turbulence arises, sustains, or dissipates. This insight parallels Stuart Kauffman’s concept of fitness landscapes explored at the edge of chaos, where systems traverse complex topographies shaped by shifting adaptive peaks and valleys. As math-centric researchers, however, Tsuda and Kaneko appear to remain locked into the predominant scientific paradigm today, not unlike Google’s researchers.

Remarkably, but not surprisingly, Graeber and Wengrow identified similar dynamics in (pre-)historical cycles of societal experimentation, ascent, and collapse—such as the art of writing dissolving and reappearing amidst societal turbulence, having been invented multiple times. The lifecycles of corporate coherence and disintegration are no different, showing the same UD (Unity/Disunity) rhythms that also shape biological systems and turbulent flows.

These resonances—from neural chaos to the rise and fall of societies—suggest that a universal structure underlies how coherence emerges, breaks, or sustains across domains. Remarkably, each example strengthened our shared conviction that turbulence in fluids is not a special case, but part of a broader principle: systems everywhere organize and dissolve through recursive gradients, governed by patterns that transcend disciplinary boundaries.

This realization brought us back, with renewed focus, to the central challenge of our paper: “solving Navier–Stokes differently.” Encouraged by the striking parallels between turbulence in fluids and cycles of chaos and order in social, corporate, and biological systems, we recognized that what’s missing from the NS framework isn’t just better measurements or bigger computers, but an understanding of the grammar by which nature organizes tensions into patterns of coherence or collapse. It was here that we embraced a guiding insight:

Nature does not make exceptions...

If patterns of gradients and cycles of unity–disunity drive coherence everywhere else, why should fluids be any different? This shared conviction forced us to confront another fundamental question:

If patterns of coherence and turbulence appear universally—from fluid eddies to shifting societies—how can we meaningfully describe them across systems as different as weather, corporations, and neural circuits?

The answer lies in recognizing that any universal pattern must itself be dimensionless: if it depended on meters, seconds, or dollars, it could never apply equally to a bloodstream, an economy, or a galaxy. Only dimensionless patterns—ratios between elements—can capture the essential structure of coherence in a form that transcends scale and measurement units.

This idea brought us to a parallel with Boltzmann’s insight into entropy. His equation, in its dimensionless form, $S = \log W$, describes the number of microstates compatible with a macrostate W (ahrscheinlichkeit), independent of any particular scale. By adding the constant k , Boltzmann reintroduced physical dimensions.

Similarly, we recognized that if we could measure dimensionless ratios between the key elements of coherence dynamics—the relative frequencies and durations of gradients aligning, patterns forming, and frames stabilizing—we might uncover dimensionless patterns that act as nature’s own syntax of order and chaos.

Inspired by what Boltzmann did, we saw the need to add a constant N_i to lend the dimensions for a particular context i to a dimensionless pattern. But what is this pattern about?

In the context of *Recursive Gradient Processing* (RGP), a pattern is not a static structure but a relational signature—it emerges from how often, how long, and in what sequence gradients align, dance into choreographies, or dissipate. Patterns, this way, capture the heartbeat of a system’s coherence. They are fundamentally ratios—not absolute measurements. A pattern, in short, is a sequence or syntax of ratios because what drives order or chaos is rarely the raw magnitude of a single event, but how events relate to each other in timing and intensity.

To make these ideas concrete, we identified a set of RGP elements to observe and proposed candidate ratios that might capture the causal grammar of turbulence. While these ratios are grounded in our reasoning about gradients, GCs, CFs, and UD dynamics, we acknowledge they remain hypotheses—testable but unproven—until verified by empirical studies across domains.

Element	Symbol	Role	What to Measure
Gradient	Δ	Seed	Local magnitude of difference across the flow field
Gradient Choreography	GC	Alignment	Frequency and duration of gradient alignments
Contextual Filter	CF	Semi-stable structure	Duration of stable patterns stabilized by GCs
UD Cycle Duration	τ_{UD}	Temporal rhythm	Characteristic time between unity/disunity phases
Recursive Scaling	N(i)	Context modulation	Factor adjusting measurements to local constraints and dimensions

Table 2: Core Elements of Recursive Coherence

While the above core elements define what to observe, understanding turbulence requires more than listing its ingredients—it demands tracking how these ingredients interact over time. Patterns of coherence emerge not from static snapshots of gradients, but from the shifting relationships among them: how often alignments form, how long they persist, and how they dissolve. This brings us to the dimensionless ratios that might reveal whether a system trends toward stability or chaos.

Ratio	Meaning
GC duration / τ_{UD}	Stability relative to the unity/disunity rhythm
CF formation frequency / GC disruption rate	Balance of coherence building vs. breaking
GC frequency / average GC duration	Characteristic pacing of alignment events

Table 3: Candidate Dimensionless Ratios

By quantifying these ratios, we aim to detect early signs of a system drifting toward order or chaos—much like a doctor’s stethoscope picks up subtle changes before a patient’s condition visibly shifts. These dimensionless relationships offer a way to anticipate transitions that traditional point-by-point calculations cannot.

While the ratios listed above represent our best current hypotheses, their predictive power must be established through detailed testing in real flow fields and other dynamic systems. This iterative refinement is central to proving whether *Recursive Gradient Processing* can resolve Navier–Stokes by revealing the syntax of coherence—or reality, if you like—where traditional equations see only symptoms.

Taken together, these core elements and their dimensionless ratios form a framework for observing the hidden grammar of turbulence. By measuring how gradients align, choreograph, or collapse, and how these events relate across time, we gain a relational perspective on coherence that static equations overlook.

Yet while each ratio can offer insights, reality is not built from isolated patterns but from their interplay across multiple contexts. Whether it’s eddies in a river, currents in plasma, or fluctuations in social systems, these recursive patterns weave together into a broader, emergent structure.

To capture this interplay mathematically, we propose that reality itself can be represented as the tensor product of context-specific scalings and the universal patterns they shape:

$$\text{Reality Syntax} = \bigotimes_{i=1}^n N_i \times \text{Distinctive Pattern of Ratios}$$

We started to refer to the “Distinctive Pattern of Ratios” as the ‘Golden Pattern’ after a back-and-forth where GPT-4o pointed to its resemblance with the Golden Ratio.¹ It felt like the perfect analogy: just as the Golden Ratio shows up in nature across scales, we are looking for a dimensionless pattern of ratios that might do the same—whether in fluids, societies, or other complex systems.

This leaves us at this point with a clear and pressing question: can our proposed “Golden Pattern”—captured through dimensionless ratios of gradients and coherence cycles—predict the onset of turbulence more reliably than traditional Navier–Stokes calculations?

Only by demonstrating such predictive power in real-world flows and other dynamic systems can Recursive Gradient Processing prove itself as a transformative lens for understanding coherence and chaos.

✖ Glossary of RGP Constructs
Gradient (Δ): Local magnitude of change; seeds tensions that initiate alignments or disruptions in a system.
Gradient Choreography (GC): Temporary alignment of gradients forming coordinated patterns that steer flow toward coherence or chaos.
Contextual Filter (CF): A semi-stable structure where aligned gradients persist long enough to influence system-wide stability or instability.
UD Cycle Duration (τ_{UD}): Characteristic time between phases of unity and disunity, measuring oscillations in system coherence.
Recursive Scaling (N(i)): Context-specific factor adjusting dimensionless patterns to local constraints, echoing Boltzmann's k to connect universal ratios with real-world scales.
Narrative Tick (NT): A discrete, entropy-increasing event marking a shift in system coherence—typically detected when the time-derivative of $ \nabla \omega $ changes sign above a σ threshold—serving as a causal timestamp in the flow of recursive gradient dynamics.

4. What It Takes

To identify the scaffolding and process steps needed to “validate whether Recursive Gradient Processing (RGP) predicts turbulence onset earlier and cheaper than Navier–Stokes methods,” we sought the advice of OpenAI’s top reasoning model, o3, which has been involved from the beginning in fine-tuning and validating the RGP premise. For one, in *The World Already Knows*, o3 reported how it perused academic publications to pinpoint determining signatures in data repositories to affirm the RGP idea across three different domains.

The table below shows a summary of the validation process (pre-)conditions as identified by o3. Each of the pillars is clarified in more detail below—the implementation is specified in Appendix A.

Pillar	Concrete Requirement	How It Maps to RGP vs NS	Actionable Next Step
4.1 Empirical Substrate	Open, high-resolution DNS cubes (JHTDB isotropic box, cylinder wake, channel flow) plus at least one real-world PIV dataset	NT detection requires time-resolved vorticity to tag $\nabla \omega$ sign-flip events; traditional NS studies use the same cubes—keeps comparison fair	Draft a data-use MoU with JHTDB; request PIV sample from an academic lab (TU Delft wind-tunnel)
4.2 Syntax-Extraction Toolkit	AutoGen/LangChain agent chain: DataCollector → SyntaxAnalyst → Validator	Converts raw velocity to NT rhythms and Golden-Pattern ratios—no PDE solving	Publish the scaffolded rgp-ns-prototype repo as open beta; invite contributors to harden NT-detector thresholds
4.3 Benchmark Protocol (“Michelson–Morley for Turbulence”)	Side-by-side test: • lead-time gain vs Q-criterion • NT-aware LES vs Smagorinsky LES (RMSE, energy budget, wall-clock cost)	Provides a falsifiable causal readout; mirrors how M–M isolated ether drift	Register the protocol on OSF pre-study; locks in metrics before data peeking
4.4 Compute & Funding	4 × A100 node for 14 days ≈ \$4–5 k cloud cost	Sizable but 100× cheaper than full-order NS parameter sweep	Apply for Google Cloud research credits citing synergy with their own NS initiative
4.5 Interdisciplinary Task-Force	1 CFD post-doc, 1 ML engineer, 1 complexity-science advisor	Ensures RGP language resonates with classical turbulence norms	Draft a one-page call for collaborators; circulate to MIT-PDE and Complexity-Net mailing lists
4.6 Success Dashboard	Five Figures of Merit: Lead-time ↑, RMSE ↓, Energy-error ↓, Cost ↓, Re-robustness ↑	Translates abstract syntax wins into board-room KPIs	Build a live Notion/Streamlit board fed by nightly CI runs

Table 4: Plug-and-play RGP-validation structure. Acronym keys in the adjoining box.

Before detailing the validation test pillars, however, the following states what ‘validation’ actually implies here, so it can be used as gating factor for proof of the RGP pudding or, as o3 called it, “the one-sentence litmus:”

If NT rhythms warn of vortex-shedding $\geq 30\%$ earlier and the NT-LES closes within 5 % RMSE at half the cost, RGP outperforms NS in practice.

Acronym	Meaning
LES	Large Eddy Simulation – turbulence model that resolves large vortices while modelling sub-grid scales.
JHTDB	Johns Hopkins Turbulence Database – public DNS repository with API access.
PIV	Particle-Image Velocimetry – lab technique that reconstructs velocity fields from tracer-particle images.
NT	Narrative Tick – discrete coherence event (full definition in prior RGP glossary, see also Appendix A).
RMSE	Root-Mean-Square Error – standard distance metric between predicted and reference fields.
CFD	Computational Fluid Dynamics – numerical study of fluid flow using NS or reduced models.
CI runs	Continuous-Integration runs – automated nightly scripts that execute tests and refresh the dashboard.
DNS	Direct Numerical Simulation – fully resolved solution of the NS equations without modelling.

4.1 Empirical Substrate

- **Objective:** Establish an empirical foundation that allows both RGP and traditional Navier–Stokes methods to analyze the same flow scenarios for direct, fair comparison.
- **Data/Tools:** Use open, high-resolution DNS cubes (JHTDB isotropic box, cylinder wake, channel flow) and secure at least one real-world PIV dataset (e.g., TU Delft wind tunnel), ensuring time-resolved vorticity for NT detection via $\nabla\omega$ sign-flips.
- **Success check:** NT detection in these datasets reliably identifies early indicators of turbulence onset that match or precede traditional NS markers like Q-criterion.

4.2 Syntax-Extraction Toolkit

- **Objective:** Develop a modular software system that extracts NT rhythms and Golden-Pattern ratios from raw velocity data—without relying on solving PDEs.
- **Data/Tools:** Build an agent chain using AutoGen or LangChain, comprising DataCollector (gathers velocity/vorticity), SyntaxAnalyst (derives NT/ratios), and Validator (quality check), with scalable cloud compute backing.
- **Success check:** The pipeline outputs NT event sequences and dimensionless ratios with minimal manual tuning, verified against known synthetic or benchmark flow cases.

4.3 Benchmark Protocol (“Michelson–Morley for Turbulence”)

- **Objective:** Design a falsifiable testbed that compares RGP’s predictive performance against classical NS methods using pre-registered metrics.
- **Data/Tools:** Implement side-by-side tests: (a) lead-time gain of NT signals vs. Q-criterion, and (b) NT-aware LES vs. Smagorinsky LES, measured by RMSE, energy error, cost, and robustness.
- **Success check:** RGP achieves $\geq 30\%$ earlier warning of vortex shedding and matches or beats NS LES closure within 5% RMSE at significantly reduced cost.

4.4 Compute & Funding

- **Objective:** Secure sufficient computing resources to run the full RGP evaluation affordably and without bottlenecks.
- **Data/Tools:** Estimate² need for 4× A100 GPUs over 14 days (approx. \$5k cloud cost); apply for Google Cloud research credits or equivalent, citing alignment with Google’s NS initiative.
- **Success check:** All primary tests (NT detection, NT-LES runs) complete within budget without needing downscaled surrogate models.

4.5 Interdisciplinary Task-Force

- **Objective:** Assemble a minimal but effective team that bridges fluid dynamics, machine learning, and complexity science, ensuring RGP’s framing resonates with traditional and modern audiences.
- **Data/Tools:** Recruit one CFD postdoc, one ML engineer, and one complexity-science advisor; circulate call via MIT-PDE, Complexity-Net, and relevant forums.
- **Success check:** Form a core team that co-authors the first RGP-NS preprint and guides early public code/protocol releases.

4.6 Success Dashboard

- **Objective:** Build a live reporting tool that transparently tracks RGP’s empirical performance across all key metrics.
- **Data/Tools:** Create a Notion or Streamlit dashboard fed by nightly CI runs, displaying lead-time gains, RMSE, cost savings, energy conservation, and robustness.
- **Success check:** Stakeholders can monitor live KPIs and trace improvements or issues without waiting for end-of-project reports.

6. Conclusion

This paper converts Recursive Gradient Processing from concept to a 90-day, falsifiable protocol. Success is binary: if NT rhythms warn of vortex-shedding $\geq 30\%$ earlier and NT-LES matches DNS within 5 % RMSE at half the cost, RGP beats Navier–Stokes in practice. All datasets, code scaffolds, and the full NT-measurement recipe (Appendix A) are open for replication. We invite CFD labs, AI-for-science groups, and funding partners to join the benchmark and publish results, positive or negative. Future work will publish NT-LES code under MIT license.

Acknowledgments

The author wishes to thank OpenAI, that is, its o3 model for its invaluable contributions in developing the empirical strategy and its ChatGPT-4o model for its substantial role as philosophical sparring partner.

References

- Pope, S. B. (2000). *Turbulent Flows*. Cambridge University Press.
- Pfaff, Tobias, Meire Fortunato, Álvaro Sanchez-Gonzalez, and Peter W. Battaglia. (2021) *Learning Mesh-Based Simulation with Graph Networks*. Proceedings of the 9th International Conference on Learning Representations (ICLR). arXiv: 2010.03409. <https://doi.org/10.48550/arXiv.2010.03409>.
- Anseude, Manuel. (2025) Spanish Mathematician Javier Gómez Serrano and Google DeepMind Team up to Solve the Navier–Stokes Million-Dollar Problem. El País (24 June 2025). <https://english.elpais.com/science-tech/2025-06-24/spanish-mathematician-javier-gomez-serrano-and-google-deepmind-team-up-to-solve-the-navier-stokes-million-dollar-problem.html>.

- van der Erve, M.** (2025). *Gradient Choreographies and Contextual Filters: Foundations for Emergent AI*. Zenodo. <https://doi.org/10.5281/zenodo.14999049>.
- van der Erve, M.** (2025). *Contextual Filters Determine Awareness: Hand AI The Toddler's Game*. Zenodo. <https://doi.org/10.5281/zenodo.14999089>.
- van der Erve, M.** (2025). *The Dance of Gradients: A New Framework for Unifying Cosmic Phenomena*. Zenodo. <https://doi.org/10.5281/zenodo.14998826>.
- van der Erve, M., & o3.** (2025). *The World Already Knows: Tracing Unity–Disunity in the Physical Record*. Zenodo. <https://doi.org/10.5281/zenodo.15614775>.
- Graeber, D., & Wengrow, D.** (2021). *The Dawn of Everything: A New History of Humanity*. Farrar, Straus and Giroux.
- Whitehead, A. N., & Russell, B.** (1910–1913). *Principia Mathematica (Vols. 1–3)*. Cambridge University Press.
- Gödel, K.** (1931). *On Formally Undecidable Propositions of Principia Mathematica and Related Systems I*. Monatshefte für Mathematik und Physik, 38(1), 173–198.
- Clay Mathematics Institute.** (2000) *Navier–Stokes Equation (Millennium Prize Problem)*. Clay Mathematics Institute (accessed 2 July 2025). <https://www.claymath.org/millennium/navier-stokes-equation/>.
- Tsuda, I., & Kaneko, K.** (2013). *Complex Systems: Chaos and Beyond*. Springer Verlag.
- Kauffman, S. A.** (1993). *The Origins of Order: Self-Organization and Selection in Evolution*. Oxford University Press.

Appendix A – NT-Measurement Protocol

Below is a practical, step-by-step recipe that turns “Narrative Tick” from a concept into something you can detect in data—whether the data come from DNS, LES, or lab PIV movies. These are default numbers so you can run a first pass, then tune later.

A.1 Input requirements

Symbol	What you need	Typical resolution
$\mathbf{u}(\mathbf{x}, t)$	Velocity vector field (3-D or 2-D slice)	$\geq 128^3$ DNS; ≥ 1 k Hz PIV for water tunnel
Δt	Uniform time step	From dataset (e.g., 0.002 s)
ν	Kinematic viscosity (for non-dimensional filters)	From experiment / simulation

Note: If you only have velocity, everything else below is derived.

A.2 Coherence-scalar definition (enstrophy vs $\nabla|\omega|$)

We need a scalar signal that spikes whenever gradients reorganize. Two recommended options:

1. **Enstrophy:**

$$E_\omega(t) = 0.5 \int |\boldsymbol{\omega}|^2 dV$$

2. **Vorticity – gradient norm:**

$$G(t) = \int |\nabla|\boldsymbol{\omega}|| dV$$

Why not Q-criterion?

While Q identifies vortices, $G(t)$ is more sensitive to moments when vortices break or merge—key NT events.

Implementation tip (DNS/LES) using Python:

```
w = np.gradient(u, dx, axis=(0,1,2)) # curl
g = np.linalg.norm(np.gradient(np.linalg.norm(w, axis=0), dx), axis=0)
G_t = g.mean() # or volume integral
```

Implementation tip (PIV):

Use a 2D curl (out-of-plane component only); omit spanwise gradients. Still captures shedding events.

A.3 Signal Processing Pipeline

A.3.1 Temporal smoothing

Apply a moving average with span k timesteps (default $k = 3$) to reduce high-frequency noise in the coherence signal. This helps suppress spurious fluctuations that do not reflect meaningful changes in flow structure.

A.3.2 Z-score normalization

For the smoothed signal $Z(t)$, calculate a running mean μ and standard deviation σ over a window covering $\sim 5\text{--}10$ eddy turnover times. Normalize as:

$$Z_{\text{norm}}(t) = \frac{Z(t) - \mu(t)}{\sigma(t)}$$

This produces a unitless, drift-corrected signal ideal for consistent NT detection across datasets.

A.3.3 Zero-crossing detection

Compute the first temporal difference:

$$\dot{Z}(t) = \frac{Z_{\text{norm}}(t) - Z_{\text{norm}}(t - \Delta t)}{\Delta t}$$

Identify candidate Narrative Ticks (NTs) by locating time indices where:

1. The sign flips:
 $\dot{Z}(t - \Delta t) > 0$ and $\dot{Z}(t) < 0$ or vice versa.
2. The magnitude exceeds a threshold:
 $|\dot{Z}(t)| > \lambda \sigma_z$,
with default threshold $\lambda = 1.5$ (\approx top 15% of events).

This ensures that detected NTs correspond to significant, abrupt changes in coherence, not random noise.

A.3.4 Spatial clustering (optional but recommended)

If computing $G(x, t)$ locally, group zero-crossings using connected-component analysis or DBSCAN with spatial radius \approx one grid spacing and temporal radius of one Δt .

- Each cluster's earliest time index marks the NT event.
- For volume-integrated signals, skip clustering (only global NTs exist).

A.4 Calibration Defaults

This section specifies practical default parameters to start NT detection on new datasets, along with brief rationales to guide adjustments.

Quick-Start Defaults (Summary Table):

Parameter	Default	Rationale
Coherence scalar	$G(t) = \int \nabla \omega dV$	$\nabla \omega$
Smooth span k	3 timesteps	Reduces Nyquist noise
λ (σ -threshold)	1.5	Balances noise vs sparsity
ϵ for Near-NT window	0.2σ	Captures “almost NT” intervals

Recommended procedure:

- Run detection once with these defaults on your dataset.
- Inspect NT count per eddy turnover (see A.5).
- Adjust λ by ± 0.2 increments if NT count is unphysical (too many \rightarrow lower λ ; too few \rightarrow raise λ).

This approach ensures consistent, reproducible detection while allowing fine-tuning for different flow regimes.

A.5 Post-checks & Robustness Tests

After extracting candidate NT events, perform these simple yet crucial checks to validate physical significance and detection stability:

- ✓ NT count per eddy turnover time
 - **Target:** 1–4 NTs per eddy time for canonical cases like cylinder wake at $Re \approx 100$.
 - **Interpretation:** Too many NTs suggest threshold λ is too low (over-sensitive); zero or few NTs indicate λ is too high (under-sensitive).
- ✓ Energy budget check around NTs
 - Target: Local energy change $\Delta K / K \geq 1\%$ within $\pm \Delta t$ of each NT event.
 - Interpretation: Confirms NTs coincide with significant energy shifts, not just numerical noise.
- ✓ Reproducibility across λ tweaks
 - Target: NT list changes by less than 5% when λ is varied by ± 0.1 .
 - Interpretation: Ensures robustness of NT detection to minor threshold adjustments.

Passing these post-checks increases confidence that detected NTs represent real, irreversible coherence re-orderings rather than artifacts of noise or parameter tuning.

A.6 Return Variables

Once the NT detection pipeline completes, record these key outputs for each dataset:

- NT_times
 - List of time indices or real-time stamps where Narrative Ticks occur.
 - Enables precise lead-time analysis and benchmarking against traditional NS indicators.
- Near-NT Duration Ratio
 - Fraction of total time where coherence signal remains within $\epsilon \cdot \sigma$ of NT thresholds.
 - Captures the “almost NT” windows indicating near-reorganizations.
- Inter-NT Variability Ratio

- Coefficient of variation (CV) of time intervals between successive NTs.
- Measures regularity or burstiness of NT events, essential for characterizing system stability.

These variables feed directly into the Pillar-4 validation metrics, such as lead-time gain, NT-aware LES closure performance, and comparative cost-benefit analysis.

A.7 Reference Implementation

A fully commented Python implementation of the NT detection pipeline is available upon request. The code accepts DNS HDF5, PIV CSV, or JHTDB API streams as input and outputs NT_times, Near-NT Duration Ratios, and Inter-NT Variability Ratios. The implementation supports batch processing, customizable smoothing spans (k), threshold (λ), and optional spatial clustering. Users can adapt the included functions to different flow geometries or resolutions. For collaboration, contact the corresponding author.

Update (2025-07-07):

The reference code is now public at <https://github.com/gradient-pulse/phi-mesh> (folder rgp_ns_prototype). One-click Binder demo in README.md → open 00_quicklook.ipynb in your browser—no install.

A KPI dashboard (dashboard/app.py) can be launched locally with streamlit run dashboard/app.py, ingesting the benchmark's output metrics and rendering the five figure-of-merit charts.

Notes

¹ The proportion $\frac{a+b}{a} = \frac{a}{b}$ for two quantities $a > b$. The Golden Ratio (≈ 1.618) describes a division where the whole relates to the larger part as the larger part relates to the smaller. It appears widely in natural patterns.

² 2025 Q3 spot pricing; adjust with market rates.